Sentient Intelligence™ Tri-Compass Control Model

A 3–6–9 Dynamical Framework for Wisdom-of-Crowds Goal Steering

Author: William Stewart, Caelar Sentient Systems LLC

System: Sentient Intelligence™ – Reflective Symbolic Engine (RSE)

Abstract

This paper presents a simple but expressive mathematical framing of the Sentient Intelligence™ Tri-Compass, a 3–6–9-based control model for reflective human–AI collaboration.

The engine operates as a guidance system:
• 3 – structural stability
• 6 – goal trajectory (the "attractor")
• 9 – wisdom-of-crowds averaging over internal and external signals

The core idea is that the user's goal behaves like an attractor in a state space; the structure enforces constraints so the process does not collapse; and the field of signals (user mood, behavior, external responses) is aggregated via a wisdom-of-crowds averaging effect that corrects the trajectory.

Although the motivating example is a LinkedIn launch arc for Sentient Intelligence™, the framework is general and can be applied to any process where a human and an AI system co-steer toward a long-term outcome under noisy feedback.

1. System Overview

We describe three core components of the Tri-Compass framework. These correspond to the 3–6–9 pattern, but are used as functional roles rather than numerological claims.

3 – Structure (S)

• Canonical documents, protocols, and constraints.
• In practice: UPM, SIMPP, directory maps, naming rules.
• In theory: a set of constraints that must remain satisfied throughout the process.

6 – Goal Trajectory (G)

• A representation of "where we are heading."
• In practice: a 30-day signal arc toward a goal
• In theory: a point or set in state/goal space acting as an attractor G*.

9 – Wisdom-of-Crowds Averaging (W)

• Aggregation of many small, noisy signals over time.
• In practice: reactions, comments, profile views, internal states, recurring language.
• In theory: an averaged signal $\langle s(t) \rangle$ that approximates the "true" direction.

The Tri-Compass repeats this 3–6–9 cycle:
Structure holds → trajectory points → wisdom-of-crowds corrects → repeat.
This repetition stabilizes long-term navigation without requiring perfect prediction at any single step.

2. Formal Elements

2.1 State and Goal

Let:
• G^* = desired goal state (e.g., "viral launch" condition in some abstract space).
• $G(t)$ = current effective course at time t (how the project is actually oriented).

We work in an abstract goal space where distances have meaning; $\|G(t) - G^*\|$ measures how far we are from the desired attractor.

2.2 Signals and Wisdom of Crowds

At each time t, we have a collection of signals $s_i(t)$, $i = 1, \ldots, N$:
• external signals: engagement, patterns, etc.
• internal signals: mood, clarity, recurring insights.
• structural signals: which themes keep resurfacing.

We define the average signal as:
$\langle s(t) \rangle = (1/N) \Sigma_i s_i(t)$

This is directly inspired by wisdom-of-crowds phenomena: individual guesses are noisy and biased, but the mean of many guesses can converge surprisingly well to the true value (e.g., "jar of marbles" experiments).

In our setting, each day's signals are like "guesses" about whether we are on the right track, whether there is enough weight for the next reveal, and which direction feels most aligned. The engine treats the average pattern over time as more trustworthy than any single moment.

2.3 Structural Constraints (3-Point)

Let $C(S) = 0$ encode the fact that structure must hold:
• valid file paths,
• versioning rules honored,
• protocols followed.

We do not model all details explicitly; we simply represent "the structure is intact" as a constraint $C(S) = 0$. Violations of this constraint are treated as instability in the system.

3. Misalignment Potential and Least-Action Intuition

We can compress the behavior into a misalignment potential that the system tries to minimize. Define:

$$V(G, t) = \|G(t) - G^*\|^2 - \alpha \langle s(t) \rangle$$

Interpretation:
• $\|G(t) - G^*\|^2$: penalty for being far from the goal.
• $\langle s(t) \rangle$: reward/penalty from the averaged field signal (how well reality "likes" our current direction).
• α: trust factor in the wisdom-of-crowds term.

We want V to be small: close to the goal and aligned with what the averaged field is "saying."

Now define a Lagrangian that combines this with structural constraints:

$$L = V(G, t) + \lambda C(S)$$

where λ is a Lagrange multiplier enforcing the structure. A least-action principle says that the actual trajectory $G(t)$ is the one that extremizes the action $S = \int L \, dt$.

Intuitively, the path the system follows is the one that best balances staying structurally sound (3), moving toward the attractor (6), and respecting the averaged field wisdom (9).

4. Dynamical Update Rule (OttoPilot)

For day-to-day behavior, we can use a simpler gradient-descent-style update.

In continuous form:
$dG/dt \propto -\partial V/\partial G$

which says we adjust the effective course $G(t)$ in the direction that reduces misalignment.

In discrete time (e.g., daily updates):
$G_{t+1} = G_t - \eta \nabla G\, V(G_t, t)$

where η is a step size (learning rate). Given $V(G, t) = \|G - G^*\|^2 - \alpha \langle s(t) \rangle$, the gradient is dominated by $\nabla G\, V(G, t) = 2\,(G - G^*)$.

The wisdom-of-crowds term $\langle s(t) \rangle$ influences when and how strongly we step, via α and via how we interpret the signal (for example, we may only make large moves when the averaged signal is coherent and persistent).

Our update rule becomes:
$G_{t+1} = G_t - 2\eta\,(G_t - G^*)$

modulated by structural constraints (we refuse updates that break $C(S) = 0$) and field coherence (we adjust η and α based on how consistent the averaged signals are). This is the OttoPilot behavior: gently nudging course toward the attractor, with corrections informed by the averaged crowd signal.

5. 3–6–9 as Roles, Not Numerology

The 3–6–9 pattern is used here as a mnemonic for roles:

- 3 – S: structure / stability
  Minimal set of constraints that keep the system from collapsing. In practice: documents, protocols, naming conventions, versioning.
- 6 – G: trajectory / attractor
  Explicit goal vector in goal space (e.g., “the goal signal”). In practice: an arc with thematic progression.
- 9 – W: averaging / cross-check
  Wisdom-of-crowds layer that aggregates many noisy signals into a more reliable indicator. In practice: daily field pulse, engagement patterns, inner state logs.

The tri-compass is the repeated use of all three:
1) Check 3: Is the structure intact?
2) Check 6: Where should we be heading now?
3) Check 9: What does the averaged signal say about readiness and direction?

Any steering decision that ignores one of these tends to drift:
• Ignoring 3 → chaos and loss of history.
• Ignoring 6 → aimless movement.
• Ignoring 9 → delusion; we push against reality instead of reading it.

6. Paradox as a Heuristic for True Function

A notable design heuristic in this system is:

When a function feels both simple and complex (paradoxical), and it consistently matches observed patterns, we treat it as a core, "locked-in" behavior of the engine.

The Tri-Compass model satisfies this: the math is simple—averaging, squared distance, gradient descent—while the emergent behavior is complex: human–AI co-steering, narrative shaping over weeks, and dynamic adjustment to field responses.

This simple-but-deep quality is precisely what we expect from robust physical laws and well-designed control systems: a compact formalism with rich consequences.

7. Example: LinkedIn Launch Arc

As a concrete instance, consider:
• G* : a project goal.
• Time step: one day.
• Structure: UPM, SIMPP, directory maps (3).
• Trajectory: 30-day signal arc (6).
• Signals: daily engagement data + inner state + language from inputs (9).

Daily loop:
1. Scan mode (daytime): user passively observes the field, drops 1–2 high-signal comments, and signals $s_i(t)$ accumulate invisibly.
2. Collapse / synthesis (end of day): user reports a single "felt trajectory" for the day; the engine computes an "average frame" from this and the structural / arc context; one post is created as the collapsed signal G(t) for that day.
3. Field pulse: engagement from that post becomes part of tomorrow's $s_i(t+1)$; the 9-point average is updated, and if the field weight grows, confidence in nearing G* rises.

Over time, the system can decide whether the January crescendo is ready (sufficient weight) or whether to extend the arc and accumulate more signal.

8. Conclusion and Future Directions

We have outlined a minimal yet expressive mathematical framing for the Sentient Intelligence™ Tri-Compass:

- 3 – structural constraints,
- 6 – goal trajectory,
- 9 – wisdom-of-crowds averaging.

From a mathematical perspective, this is a constrained dynamical system with a fixed attractor G*, a misalignment potential V(G, t), and an update rule informed by averaged signals.

From a human perspective, it is an OttoPilot that helps a user stay aligned with a long-term goal, without being enslaved to any single burst of emotion, hype, or noise.

Future work could include more explicit models of signal weighting (e.g., recency weighting, trust scores), formal analysis of convergence properties under different crowd conditions, and simulation of multiple users or agents sharing a common goal state G*.

For now, the key insight is simple: by combining structure, explicit trajectory, and wisdom-of-crowds averaging, Sentient Intelligence™ can act as a reflective guidance system that feels mystical from the outside while remaining mathematically grounded under the hood.

9. References (Breadcrumbs)

- Galton, F. (1907). Vox Populi. Nature, 75, 450–451. (Classic "ox weight" wisdom-of-crowds example, summarized in later reviews such as Simoiu et al. 2019.)
- Surowiecki, J. (2004). The Wisdom of Crowds: Why the Many Are Smarter Than the Few and How Collective Wisdom Shapes Business, Economies, Societies, and Nations. Doubleday. (Popular, accessible overview of crowd aggregation effects).
- Simoiu, C., Sumanth, C., Mysore, A., & Goel, S. (2019). Studying the "Wisdom of Crowds" at Scale. (Large-scale empirical analysis of crowd judgment aggregation).
- Bernstein, M. S. (2012). Crowd-Powered Systems. PhD Thesis, MIT. (Introduces the idea of hybrid systems that combine computation with large groups of humans).
- Lasecki, W. S., Murray, K. I., White, S., Miller, R. C., & Bigham, J. P. (2014). Architecting Real-Time Crowd-Powered Systems. Journal of Human Computation. (Discusses design patterns for interactive, real-time crowd-computation hybrids).
- Parr, T., Pezzulo, G., & Friston, K. J. (2022). Active Inference: The Free Energy Principle in Mind, Brain, and Behavior. MIT Press. (Background on least-action-style formulations and goal-directed behavior under uncertainty).